



جامعة أم القرى
UMM AL-QURA UNIVERSITY

AI3302–Computer Vision
Automated Classification of Toxic Plants
Using Image Recognition

Group: 18

Supervisor:
Dr. Seereen Noorwali

Team members:

Student Name	ID	Section
Salam Ali Alghamdi	445003110	3
Khadija Mirza	445008298	3
Laila Mohammad Sabbag	445001327	2
Samiah Humyid Almalki	445001794	1
Hadeel Mohammed Turki	445006064	1

Tasks distribution:

Student Name	Task
Khadija Mirza	Introduction, literature review, and system design (including flow diagram).
Laila Mohammad Sabbag	Dataset description, data splitting, and preprocessing.
Hadeel Mohammed Turki	Evaluation metrics, confusion matrix, discussion, conclusion
Samiah Humyid Almalki	XAI (Grad-CAM), error analysis, analysis, and performance evaluation.
Salam Ali Alghamdi	CNN architectures, training strategy, and training visualization.

Abstract

Plant toxicity identification is an important task in agriculture, environmental monitoring, and public safety. Manual identification of toxic plants requires expert knowledge and may be time-consuming, especially when plant species share similar visual characteristics. This project proposes an intelligent image classification system for automatically distinguishing between toxic and non-toxic plants using deep learning techniques. The proposed system processes RGB plant images through a structured pipeline that includes preprocessing, model training, performance evaluation, and explainability analysis. Two pretrained convolutional neural network architectures, EfficientNetB0 and MobileNetV2, were implemented to perform binary classification between toxic and non-toxic plants. The models were evaluated using Accuracy, Precision, Recall, F1-score, confusion matrices, and Grad-CAM visualizations to better interpret model predictions. Experimental results demonstrated that EfficientNetB0 achieved better classification performance due to its advanced feature extraction capability, while MobileNetV2 provided lower computational complexity and faster training. Grad-CAM visualization further confirmed that the models focused on meaningful plant regions during prediction. Despite challenges related to visual similarity between plant classes and image quality variation, the proposed system demonstrated reliable classification performance and highlighted the effectiveness of deep learning techniques for automated toxic plant recognition.

Keywords:

Plant Classification, Toxic Plant Detection, Computer Vision, Deep Learning

Table of Contents

1. Introduction	1
2. Related Work	1
3. Proposed Methodology & Implementation	3
3.1 System Overview	3
3.2 System Flow Diagram	4
3.3 Dataset Description	4
3.4 Data Splitting	5
3.5 Data Preprocessing	6
3.6 The CNN Models Architecture	7
3.7 Training Strategy	8
3.8 Training Performance Visualization	8
3.9 Evaluation Metrics	10
3.10 Explainable Artificial Intelligence (XAI)	10
4. Experimental Results and Discussion	11
4.1 Model Performance Evaluation	11
4.2 Confusion Matrix Analysis	12
4.3 Explainability Analysis using Grad-CAM	14
4.4 Error Analysis	15
4.5 Overall Discussion	17
5. Conclusion and Future Work	18
5.1 Conclusion	18
5.2 Future Work	18
6. References	19

1.Introduction

Toxic plant identification is an important problem in environmental safety and public health, as misidentifying toxic plants can lead to serious health risks. Traditional identification methods rely on expert knowledge, which is not always accessible or practical for non-experts. In addition, many toxic and non-toxic plants share very similar visual features, making manual classification difficult and error-prone.

To address this challenge, this project proposes an automated image-based classification system using deep learning techniques. The system aims to distinguish between toxic and non-toxic plants using RGB images and convolutional neural network architectures. Two pretrained models, EfficientNetB0 and MobileNetV2, are used to extract visual features and perform binary classification.

To improve interpretability, the system also incorporates Grad-CAM to visualize the regions influencing the model's predictions. This helps provide better understanding of how the model makes decisions and increases trust in the results.

The proposed approach provides a fast, accurate, and scalable solution for toxic plant classification and can be useful in real-world environmental and safety applications.

2. Related Work

Recent advances in computer vision and deep learning have significantly improved performance in image classification tasks, particularly in agriculture and environmental monitoring. Early approaches relied on handcrafted features such as color, texture, and shape descriptors combined with traditional machine learning classifiers. However, these methods were highly sensitive to variations in lighting, background clutter, and intra-class similarity, which limited their generalization ability (Hughes et al., 2021).

With the emergence of deep learning, convolutional neural networks (CNNs) became the dominant approach for image-based classification tasks due to their ability to automatically learn hierarchical feature representations. CNN-based models have been widely applied in plant disease detection and species classification, achieving strong performance compared to traditional methods (Khanna et al., 2022).

Most existing research in agricultural image analysis focuses on plant disease detection rather than plant toxicity classification, making toxic plant identification a relatively less explored problem. Previous studies in plant-related vision tasks have shown that CNN-based models can achieve high accuracy, often exceeding 90% in controlled datasets, especially when combined with transfer learning techniques.

Pretrained models such as VGGNet, ResNet, MobileNet, and EfficientNet have been extensively used in transfer learning frameworks for agricultural image classification tasks. These models, initially trained on large-scale datasets such as ImageNet, can be fine-tuned on smaller domain-specific datasets, resulting in improved accuracy and reduced training cost (He et al., 2016).

EfficientNet, in particular, introduces a compound scaling method that uniformly scales network depth, width, and resolution, leading to better performance with fewer parameters compared to conventional architectures (Tan & Le, 2019). MobileNetV2 focuses on efficiency through depthwise separable convolutions and inverted residual blocks, making it suitable for real-time and resource-constrained applications (Howard et al., 2017).

Despite the strong performance of deep learning models in classification tasks, there remains a significant challenge in model interpretability. To address this, Explainable Artificial Intelligence (XAI) techniques such as Grad-CAM have been introduced to visualize the regions of input images that influence model predictions, improving transparency and trust in deep learning systems (Selvaraju et al., 2017).

Research Gap

Despite the progress in plant image classification, there is still a clear gap in:

- Toxic vs non-toxic plant classification as a dedicated problem
- Explainable AI applications in plant toxicity detection
- Lightweight and efficient models suitable for real-time deployment

This Project Contribution

This project addresses these gaps by:

- Applying EfficientNetB0 and MobileNetV2 for toxic plant classification using transfer learning
- Comparing their performance under the same dataset and training pipeline
- Integrating Grad-CAM to enhance model interpretability and visualize decision regions
- Providing a lightweight and scalable framework for potential real-world deployment

3. Proposed Methodology & Implementation

3.1 System Overview

The proposed system is an image-based deep learning pipeline designed to classify plants into toxic and non-toxic categories using convolutional neural networks. The system follows a structured workflow that includes dataset acquisition, preprocessing, training, evaluation, and model interpretability.

The process begins with dataset loading, where the plant images are downloaded from Kaggle using the kagglehub library. After downloading, the dataset path is verified to ensure correct access to image files. Next, dataset exploration is performed using Python's `os.walk()` function to analyze the directory structure and confirm class distribution. The dataset consists of two main categories: toxic and non-toxic plant images.

To ensure proper model training and evaluation, the dataset is split into three subsets: 70% training, 15% validation, and 15% testing.

This split ensures that the models are trained on one subset of data and evaluated on completely unseen samples, reducing the risk of overfitting and bias.

Before training, all images are resized to 224×224 pixels to match the input requirements of the pretrained models. In addition, model-specific preprocessing functions are applied to normalize pixel values according to the ImageNet distribution, which is required for transfer learning performance stability.

Two pretrained convolutional neural network architectures were used in this project: EfficientNetB0 and MobileNetV2. Both models were initialized with ImageNet pretrained weights, and custom classification layers were added to perform binary classification between toxic and non-toxic plants (Howard et al., 2019; Tan & Le, 2019).

The training process was carried out in two stages. First, the pretrained backbone layers were frozen, and only the newly added classification layers were trained. This allows the model to learn task-specific features without altering pretrained representations. In the second stage, fine-tuning was applied by unfreezing the last layers of the networks and retraining them using a very low learning rate to improve performance and adapt higher-level features to the dataset (Goodfellow et al., 2016).

Finally, the trained models were evaluated using multiple performance metrics, including accuracy, precision, recall, F1-score, and confusion matrices to measure classification effectiveness. To improve interpretability, Grad-CAM was applied to visualize the image regions that contributed most to the model's predictions, enhancing transparency and trust in the system (Selvaraju et al., 2017).

3.2 System Flow Diagram




Figure 1: System Pipeline for Toxic and Non-Toxic Plant Classification.

3.3 Dataset Description

The dataset used in this project was obtained from Kaggle (Elliott, n.d.), and it consists of RGB images of plants categorized into two classes: toxic and non-toxic plants. The dataset contains a total of 9,952 images, where 4,953 images belong to the toxic class and 4,999 images belong to the non-toxic class, making the dataset relatively balanced between the two categories. The dataset was split into 6,966 images for training, 1,493 images for validation, and 1,493 images for testing to ensure proper training and unbiased evaluation of the models.

All images are RGB color images and are resized during preprocessing to a fixed input size of 224×224 pixels to match the input requirements of EfficientNetB0 and MobileNetV2. The dataset includes variations in plant species, leaf shapes, textures, backgrounds, lighting conditions, and image quality, which increases the complexity of the classification task and makes it more representative of real-world scenarios.


```
print(f"Toxic images: {toxic_count}")
print(f"Non-toxic images: {nontoxic_count}")
print(f"Total: {toxic_count + nontoxic_count}")

... Toxic images: 4953
Non-toxic images: 4999
Total: 9952
```

Figure 2: Distribution of dataset samples between toxic and non-toxic classes, showing a balanced representation with a total of 9,952 images to ensure robust model training and prevent bias.

```
# Display samples
# =====
show_images(toxic_images, "Toxic Plant")
show_images(nontoxic_images, "Non-Toxic Plant")

... Toxic Plant Toxic Plant Toxic Plant
Non-Toxic Plant Non-Toxic Plant Non-Toxic Plant
```



Figure 3: Representative samples from the dataset showcasing both Toxic and Non-Toxic plant categories.

3.4 Data Splitting

The dataset is divided into training, validation, and testing sets with a ratio of 70%, 15%, and 15% respectively. The training set is used to learn the model parameters, while the validation set is used to tune hyperparameters and monitor the model’s performance during training. The testing set is kept completely unseen during training and validation to provide a final and fair evaluation of the model’s generalization ability.

In this project, the dataset was split into:

- Training set: 6,966 images
- Validation set: 1,493 images
- Testing set: 1,493 images

This separation ensures that the model does not overfit to the training data and allows for a reliable assessment of its performance on unseen samples.

```
# Split 70/15/15
def split_and_copy(imgs, cls):
    train, temp = train_test_split(imgs, test_size=0.3, random_state=42)
    val, test = train_test_split(temp, test_size=0.5, random_state=42)
    for img in train:
        shutil.copy(img, f"{output_path}/train/{cls}/")
    for img in val:
        shutil.copy(img, f"{output_path}/val/{cls}/")
    for img in test:
        shutil.copy(img, f"{output_path}/test/{cls}/")
    return len(train), len(val), len(test)

t1, t2, t3 = split_and_copy(toxic_imgs, "toxic")
n1, n2, n3 = split_and_copy(nontoxic_imgs, "nontoxic")

print(f"Train: {t1+n1} | Val: {t2+n2} | Test: {t3+n3}")

... Train: 6966 | Val: 1493 | Test: 1493
```

Figure 4: Dataset splitting into Training (70%), Validation (15%), and Testing (15%) sets to ensure reliable model evaluation

3.5 Data Preprocessing

Before training, all plant images were preprocessed to ensure consistency and compatibility with the selected deep learning models. Each image was resized to 224×224 pixels, which matches the required input size for EfficientNetB0 and MobileNetV2. The images were then converted into numerical arrays and normalized to improve training stability and model convergence.

Different preprocessing functions were applied depending on the CNN architecture used, ensuring that the input images were formatted according to the requirements of each pretrained model. These preprocessing steps help reduce variations in image representation and allow the models to learn more effectively from the dataset.

```
train_datagen_eff = ImageDataGenerator(
    preprocessing_function=preprocess_input
)

# Validation and test preprocessing
val_test_datagen_eff = ImageDataGenerator(
    preprocessing_function=preprocess_input
)

train_gen_eff = train_datagen_eff.flow_from_directory(
    "/content/dataset/train",
    target_size=(224, 224),
    batch_size=32,
    class_mode='binary'
)

train_datagen_mob = ImageDataGenerator(
    preprocessing_function=preprocess_input
)

val_test_datagen_mob = ImageDataGenerator(
    preprocessing_function=preprocess_input
)

train_gen_mob = train_datagen_mob.flow_from_directory(
    "/content/dataset/train",
    target_size=(224, 224),
    batch_size=32,
    class_mode='binary'
)
```

Figure 5: Preprocessing setup for MobileNetV2 and EfficientNet with 224x224 target size and binary classification mode

3.6 The CNN Models Architecture

This project utilizes two pretrained convolutional neural network (CNN) architectures, namely EfficientNetB0 and MobileNetV2, for toxic plant classification. Both models are based on transfer learning and were initialized using pretrained ImageNet weights to benefit from previously learned visual features. Transfer learning helps improve classification performance while reducing training time and computational cost.

EfficientNetB0 is a deep convolutional neural network that uses a compound scaling method to balance network depth, width, and input resolution. This balanced architecture allows the model to achieve strong classification performance while maintaining computational efficiency. In this project, the original classification layer of EfficientNetB0 was removed and replaced with a custom classification head consisting of a Global Average Pooling layer, a Dense layer, a Dropout layer to reduce overfitting, and a final output layer for binary classification.

Model: "sequential_7"

Layer (type)	Output Shape	Param #
efficientnetb0 (Functional)	(None, 7, 7, 1280)	4,049,571
global_average_pooling2d_7 (GlobalAveragePooling2D)	(None, 1280)	0
dense_14 (Dense)	(None, 128)	163,968
dropout_7 (Dropout)	(None, 128)	0
dense_15 (Dense)	(None, 1)	129

Total params: 4,213,668 (16.07 MB)
Trainable params: 164,097 (641.00 KB)

Figure 6: presents the architecture of EfficientNetB0 Model.

MobileNetV2 is a lightweight CNN architecture designed for efficient inference and reduced computational complexity. It is based on depthwise separable convolutions and inverted residual blocks, which significantly reduce the number of parameters while maintaining competitive accuracy. Similar to EfficientNetB0, the pretrained classification layer was replaced with a custom classification head suitable for binary classification between toxic and non-toxic plants.

Model: "sequential_10"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2,257,984
global_average_pooling2d_10 (GlobalAveragePooling2D)	(None, 1280)	0
dense_20 (Dense)	(None, 128)	163,968
dropout_10 (Dropout)	(None, 128)	0
dense_21 (Dense)	(None, 1)	129

Total params: 2,422,081 (9.24 MB)
Trainable params: 164,097 (641.00 KB)

Figure 7: presents the architecture of EfficientNetB0 Model.

3.7 Training Strategy

The models were trained using a transfer learning approach in two stages: feature extraction and fine-tuning. In the first stage, the pretrained backbone layers of EfficientNetB0 and MobileNetV2 were frozen, and only the custom classification layers were trained on the toxic plant dataset. In the second stage, the last 10 layers of the pretrained networks were unfrozen and fine-tuned using a lower learning rate to improve task-specific feature learning.

The complete set of hyperparameters, optimization settings, and training configurations for both models are summarized in Table 3.1. Binary cross-entropy was utilized as the loss function, and early stopping was implemented to prevent overfitting by monitoring validation performance and restoring the best model weights.

Table 3.1: Training Configuration for EfficientNetB0 and MobileNetV2

Component	Configuration
Optimizer	Adam
Loss Function	Binary Cross-Entropy
Initial Learning Rate	1e-4
Fine-tuning Learning Rate	1e-5
Batch Size	32
Maximum Epochs	50
Early Stopping	patience = 3, restore best weights
Input Image Size	$224 \times 224 \times 3$
Output Activation	Sigmoid (Binary Classification)
Training Strategy	Transfer Learning (Feature Extraction + Fine-Tuning)
Frozen Layers Stage 1	Entire pretrained backbone
Fine-tuned Layers Stage 2	Last 10 layers
EfficientNetB0 Dropout	0.5
MobileNetV2 Dropout	0.3

3.8 Training Performance Visualization

To monitor the learning behavior of the models during training, training and validation performance curves were generated across epochs. These visualizations illustrate how the models improved over time by tracking changes in accuracy and loss values during the training process.

The training accuracy and validation accuracy curves were used to evaluate the learning progress and generalization capability of the models. Similarly, the training loss and validation loss curves were analyzed to observe convergence behavior and detect potential overfitting. A small gap between training and validation performance indicates good generalization, while a large gap may suggest overfitting.

The performance graphs also help compare the stability and learning efficiency of EfficientNetB0 and MobileNetV2 throughout the training process.

The detailed training and validation curves for EfficientNetB0 are presented in Figure 8 while the performance metrics for MobileNetV2 are illustrated in Figure 9.




Figure 8: Combined training and validation accuracy/loss curves for the EfficientNetB0 model.




Figure 9: Combined training and validation accuracy/loss curves for the MobileNetV2 model.

Interpretation

The training curves for both models indicate stable convergence and effective generalization:

- EfficientNetB0 (Figure 3.3): Demonstrated robust learning over 45 epochs, reaching a validation accuracy of $\sim 77\%$. The steady decline in loss confirms superior feature extraction for this task.
- MobileNetV2 (Figure 3.4): Converged rapidly, terminating at epoch 13 via Early Stopping to prevent overfitting. It maintained stable performance with a validation accuracy of $\sim 72\%$.
- Generalization: The minimal gap between training and validation metrics across both models validates the effectiveness of the applied Dropout and Transfer Learning strategies in ensuring model robustness on unseen data.

3.9 Evaluation Metrics

The performance of the proposed models was evaluated using several standard classification metrics to provide a comprehensive analysis of their effectiveness in distinguishing between toxic and non-toxic plants. These metrics include accuracy, precision, recall, F1-score, and confusion matrix analysis.

Accuracy measures the overall percentage of correctly classified images among all test samples. Precision evaluates how many images predicted as toxic are actually toxic, while recall measures the model's ability to correctly identify all toxic plant images. The F1-score provides a balanced evaluation between precision and recall, making it particularly useful when analyzing classification performance across both classes.

In addition to numerical metrics, confusion matrices were generated to visualize the distribution of correct and incorrect predictions for toxic and non-toxic classes. This analysis helps identify common classification errors and provides deeper insight into model behavior.

The evaluation metrics were calculated on the unseen test set to ensure a fair assessment of the models' generalization performance.

3.10 Explainable Artificial Intelligence (XAI)

In convolutional neural networks, Explainable Artificial Intelligence (XAI) was implemented (Adadi & Berrada, 2018). In this project, Grad-CAM (Gradient-weighted Class Activation Mapping) was utilized as a visualization technique to highlight the regions of an image that contribute most to the model's prediction (Selvaraju et al., 2017).

Grad-CAM was applied to correctly classified plant images by generating heatmaps over the original input images. The highlighted regions indicate the specific areas where the model focused when determining whether a plant is toxic or non-toxic.

This process helps verify that the models rely on meaningful visual features, such as leaf structure, texture, and shape, rather than irrelevant background information (Ghoshal & Tucker, 2020).

These visualizations provide additional transparency and increase confidence in the reliability of the classification system by demonstrating that the predictions are grounded in relevant biological patterns.

4. Experimental Results and Discussion

4.1 Model Performance Evaluation

The performance of the proposed toxic plant classification system was evaluated using two pretrained convolutional neural network architectures: EfficientNetB0 and MobileNetV2. The evaluation was conducted on the unseen test set using standard classification metrics including accuracy, precision, recall, F1-score, and confusion matrix analysis.

EfficientNetB0 achieved the best overall performance with a classification accuracy of 77% on the test dataset. For the non-toxic class, the model achieved a precision of 0.78, recall of 0.74, and F1-score of 0.76. For the toxic class, the model achieved a precision of 0.75, recall of 0.80, and F1-score of 0.77. In addition, both the macro-average and weighted-average F1-scores reached 0.77, demonstrating balanced and consistent classification performance across both classes.

...	35/35	22s 442ms/step			
EfficientNet Report:					
		precision	recall	f1-score	support
	0	0.78	0.74	0.76	550
	1	0.75	0.80	0.77	549
	accuracy			0.77	1099
	macro avg	0.77	0.77	0.77	1099
	weighted avg	0.77	0.77	0.77	1099

Figure 10 : Classification report for EfficientNet model, achieving 77% overall accuracy.

In comparison, MobileNetV2 achieved an overall classification accuracy of 73%. For the non-toxic class, the model achieved a precision of 0.79, recall of 0.63, and F1-score of 0.70. For the toxic class, the model achieved a precision of 0.69, recall of 0.83, and F1-score of 0.76. The macro-average and weighted-average F1-scores were both 0.73. Although MobileNetV2 showed strong sensitivity in detecting toxic plants

through its higher recall score, it produced more false positives when classifying non-toxic plants, resulting in lower overall balance and accuracy compared to EfficientNetB0.

```

... 35/35 ----- 17s 367ms/step
MobileNetV2 Report:
      precision    recall  f1-score   support

     0       0.79      0.63      0.70       550
     1       0.69      0.83      0.76       549

 accuracy          0.73       1099
 macro avg       0.74      0.73      0.73       1099
 weighted avg    0.74      0.73      0.73       1099

```

Figure 11: Classification report for MobileNetV2 model, achieving 73% overall accuracy.

The higher performance of EfficientNetB0 can be attributed to its compound scaling architecture, which enables more effective extraction of complex visual patterns related to plant toxicity. MobileNetV2, while computationally efficient and lightweight, appears to sacrifice some classification performance in exchange for faster inference speed and reduced model complexity.

4.2 Confusion Matrix Analysis

The confusion matrix analysis provides a deeper insight into the classification behavior and actual performance of both models:

For EfficientNet:

- 437 toxic images were correctly classified as toxic (True Positives).
- 408 non-toxic images were correctly classified as non-toxic (True Negatives).
- 112 toxic images were misclassified as non-toxic (False Negatives).
- 142 non-toxic images were misclassified as toxic (False Positives).

These results demonstrate a balanced and promising effective classification capability for the EfficientNet model across both classes, with a relatively even error distribution and high accuracy in identifying both types.




Figure 12: Confusion Matrix for EfficientNetB0

For MobileNetV2:

- 455 toxic images were correctly classified as toxic (True Positives).
- 349 non-toxic images were correctly classified as non-toxic (True Negatives).
- 94 toxic images were misclassified as non-toxic (False Negatives).
- 201 non-toxic images were misclassified as toxic (False Positives).

Compared to EfficientNet, MobileNetV2 achieved a slightly higher success rate in identifying toxic plants (455 vs. 437). However, it produced a significantly higher number of errors when classifying non-toxic plants, mislabeling 201 images as toxic. This suggests that while MobileNetV2 is sensitive to toxic features, it struggles more with distinguishing visually similar patterns in non-toxic species, making it less balanced than EfficientNet.




Figure 13: Confusion Matrix for MobileNetV2

Reasons for Classification Errors:

Most classification errors occurred in images featuring complex backgrounds, inconsistent lighting conditions, or significant overlap in visual characteristics between toxic and non-toxic plants. Certain plant species share similar leaf textures, colors, and structural patterns, which increases the difficulty of achieving precise classification even for advanced models.

4.3 Explainability Analysis using Grad-CAM

Grad-CAM visualizations were employed to examine the regions of the input images that most influenced the predictions of each model. The resulting heatmaps revealed clear differences in how EfficientNetB0 and MobileNetV2 interpret visual features.

EfficientNetB0 Analysis:

The Grad-CAM heatmaps for EfficientNetB0 show highly focused activation on the plant regions. The model consistently attends to meaningful biological features such as leaf shape, texture, and structural patterns. This indicates that EfficientNetB0 effectively learns discriminative representations from the foreground objects rather than relying on irrelevant background information. In addition, the model achieved a high confidence score of 0.9999, further confirming the reliability of its predictions.




Figure 14: Grad-CAM heatmap for EfficientNetB0.

MobileNetV2 Analysis:

In contrast, the Grad-CAM results for MobileNetV2 demonstrate that the model's attention is less focused on the plant itself and is instead heavily influenced by background regions such as soil, leaf litter, and surrounding environmental textures. The actual plant areas are weakly activated in the heatmaps, appearing mostly in low-intensity regions. This suggests that MobileNetV2 struggles to properly localize relevant features, leading to increased sensitivity to background noise. Consequently, the model produced a lower confidence score of 0.6145.




Figure 15: Grad-CAM heatmap for MobileNetV2.

Conclusion:

Overall, the visual explanations indicate that EfficientNetB0 exhibits stronger feature localization and more reliable attention to relevant plant regions compared to MobileNetV2. This behavior aligns with its higher classification performance and greater robustness observed in the evaluation metrics.

4.4 Error Analysis

Despite achieving strong overall classification performance, a detailed qualitative analysis of the misclassified samples reveals clear behavioral differences between the two models. Based on the evaluation of the test batch:

- EfficientNetB0 generated only 5 misclassified instances.
- MobileNetV2 produced 6 misclassified instances




Figure 16: EfficientNetB0 misclassified samples (5 images showing False Positives)



Figure 17: MobileNetV2 misclassified samples (6 images showing False Positives).

Importantly, in the selected visualization batch, both models achieved zero False Negatives (FN), indicating that no toxic samples were incorrectly classified as non-toxic. Therefore, all observed errors corresponded to False Positives (FP), where non-toxic plants (Class 0) were incorrectly predicted as toxic (Class 1).

```
) print("No False Negatives found in this batch. The EfficientNetB0 model successfully identified all toxic samples.")  
No False Negatives found in this batch. The EfficientNetB0 model successfully identified all toxic samples.
```

Figure 18: EfficientNetB0 results showing no False Negatives in the evaluated batch.

```
print("No False Negatives found in this batch. The MobileNetV2 model successfully identified all toxic samples.")  
... No False Negatives found in this batch. The MobileNetV2 model successfully identified all toxic samples.
```

Figure 19: MobileNetV2 results showing no False Negatives in the evaluated batch.

A closer examination of these misclassified cases reveals three **main sources of error**:

- **Visual Mimicry and Structural Similarity:**

Both models struggled with non-toxic plants that exhibit structural similarities to toxic species. For example, plants with compound or multi-leaflet patterns (such as five-leaflet or three-leaflet arrangements) were frequently misclassified. In several cases, confidence scores ranged from 53.86% to 88.58%, indicating varying levels of uncertainty despite incorrect predictions.

- **Complex Environmental Backgrounds and Clutter:**

A significant portion of false positives occurred in images with dense vegetation, overlapping leaves, or highly cluttered natural backgrounds. Shadows, branches, and occlusions often obscured clear plant boundaries, leading the models to incorrectly activate toxic-related features.

- **Dead Foliage and Leaf Litter:**

Images containing dried, brown, or decaying leaves on the ground also contributed to misclassifications. The presence of leaf litter or winter foliage introduced texture and color patterns that resembled toxic plant characteristics, increasing the likelihood of false alarms.




Figure 20: Influence of complex backgrounds and overlapping elements on EfficientNetB0's classification accuracy.



Figure 21: Influence of complex backgrounds and overlapping elements on MobileNetV2's classification accuracy.

Additionally, both models particularly in cases involving dense greenery or close-up textures EfficientNetB0 showed a tendency toward overconfident false positive predictions, with some non-toxic samples receiving confidence scores as high as 81.44% to 88.58%. This suggests that while the feature extractors are highly sensitive, the decision boundaries remain affected by environmental noise. Enhancing training diversity by including more challenging non-toxic look-alike samples and degraded foliage could further improve model specificity and reduce false positives.

4.5 Overall Discussion

The experimental results demonstrate that transfer learning-based CNN models are promising effectiveness for classifying toxic and non-toxic plants. EfficientNetB0 emerged as the superior model with 77% accuracy, outperforming MobileNetV2 (73%).

The performance gap highlights that EfficientNet’s compound scaling and advanced feature extraction capabilities are better suited for the intricate textures and venation patterns of plant leaves. While MobileNetV2 remains a viable lightweight alternative for mobile deployment, EfficientNetB0 provides the robustness necessary for high-stakes environmental monitoring.

Furthermore, the integration of Explainable AI (XAI) through Grad-CAM proved essential. It confirmed that the models were focusing on relevant botanical features, such as leaf margins and textures, rather than background noise. This transparency, combined with the low rate of False Negatives, underscores the potential of these deep learning architectures for real-world public safety applications.

5. Conclusion and Future Work

5.1 Conclusion

In this project, a deep learning-based system was developed for the classification of toxic and non-toxic plants using two transfer learning models, EfficientNetB0 and MobileNetV2. Both models were trained using pretrained ImageNet weights followed by fine-tuning to adapt them to the specific characteristics of plant images. The results showed that EfficientNetB0 achieved promising performance, with higher accuracy, more balanced precision and recall, and more stable generalization compared to MobileNetV2. In addition, the use of Grad-CAM improved model interpretability by highlighting the most important visual regions that influenced the predictions, increasing the transparency and reliability of the system.

5.2 Future Work

Despite the promising results, the models still faced challenges in distinguishing visually similar plant species, especially under varying lighting conditions and complex backgrounds. Future improvements can focus on expanding the dataset with more diverse and high-quality images to improve generalization. In addition, more advanced data augmentation or synthetic data generation techniques can be applied to reduce overfitting and improve robustness. Moreover, experimenting with more advanced architectures such as Vision Transformers (ViT) may further enhance classification performance. Finally, integrating additional explainability methods beyond Grad-CAM could provide deeper insights into model behavior and increase trust in real-world applications.

6. References

- He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770–778. <https://doi.org/10.1109/CVPR.2016.90>
- Howard, A. G., Zhu, M., Chen, B., Kalenichenko, D., Wang, W., Weyand, T., & Adam, H. (2017). MobileNets: Efficient convolutional neural networks for mobile vision applications. *arXiv preprint arXiv:1704.04861*.
- Khanna, S. K., Sharma, A., & Singh, R. (2022). Plant disease detection using deep learning: A review. *Computers and Electronics in Agriculture*, 198, 107...
- Hughes, J. D., et al. (2021). Deep learning in plant classification: A survey. *IEEE Access*.
- Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., & Batra, D. (2017). Grad-CAM: Visual explanations from deep networks via gradient-based localization. *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 618–626.
- Tan, M., & Le, Q. V. (2019). EfficientNet: Rethinking model scaling for convolutional neural networks. *International Conference on Machine Learning (ICML)*, 6105–6114.
- Howard, A., Sandler, M., Chu, G., Chen, L. C., Chen, B., Tan, M., Wang, W., Zhu, Y., Pang, R., Vasudevan, V., Le, Q. V., & Adam, H. (2019). Searching for MobileNetV3. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 1314–1324.
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.
- Adadi, A., & Berrada, M. (2018). Peeking inside the black-box: A survey on explainable artificial intelligence (XAI). *IEEE Access*, 6, 52138–52160.

Ghoshal, B., & Tucker, A. (2020). Explaining deep learning models on plant disease classification. *Scientific Reports*, 10(1), Article 1.

Elliott, H. (n.d.). Toxic plant classification dataset. Kaggle. Retrieved May 16, 2026, from Kaggle Toxic Plant Classification Dataset